

DOTE 6635: Artificial Intelligence for Business Research

# Posttraining

Renyu (Philip) Zhang

1

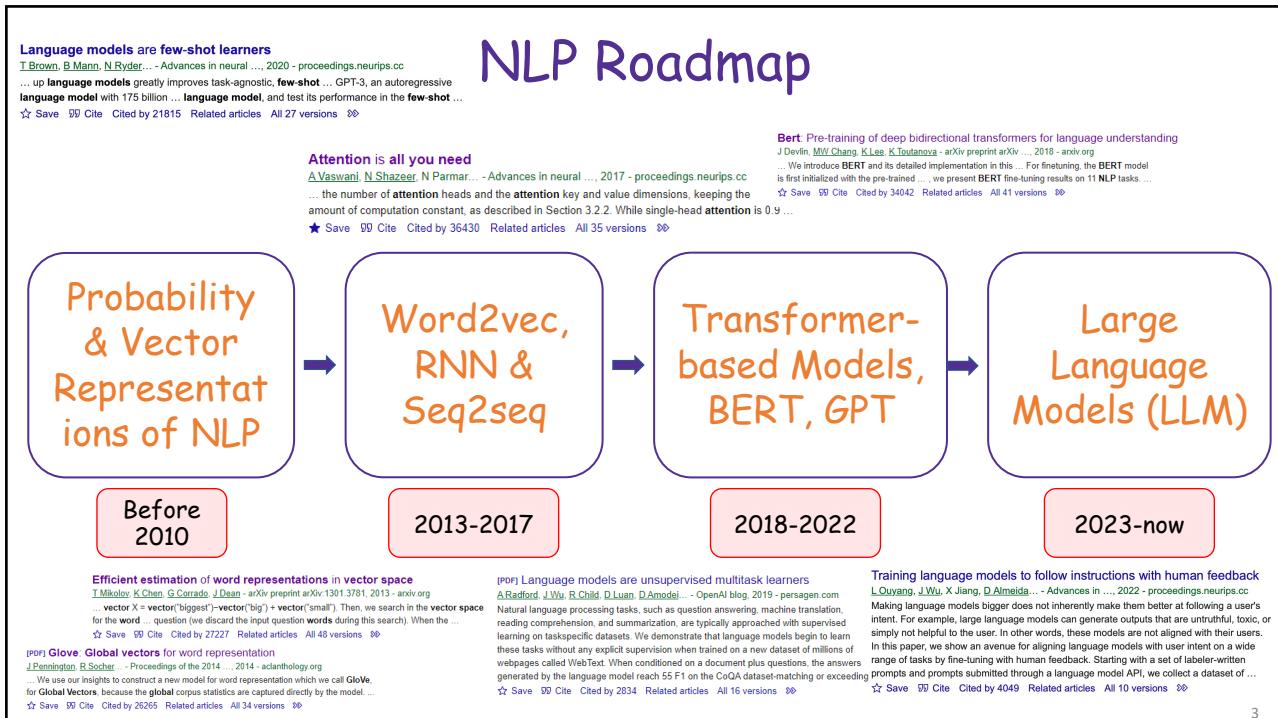
## Agenda

- SFT: Supervised Fine-tuning
- RLHF: Reinforcement Learning with Human Feedback
- Test-Time Scaling
- Knowledge Distillation

2

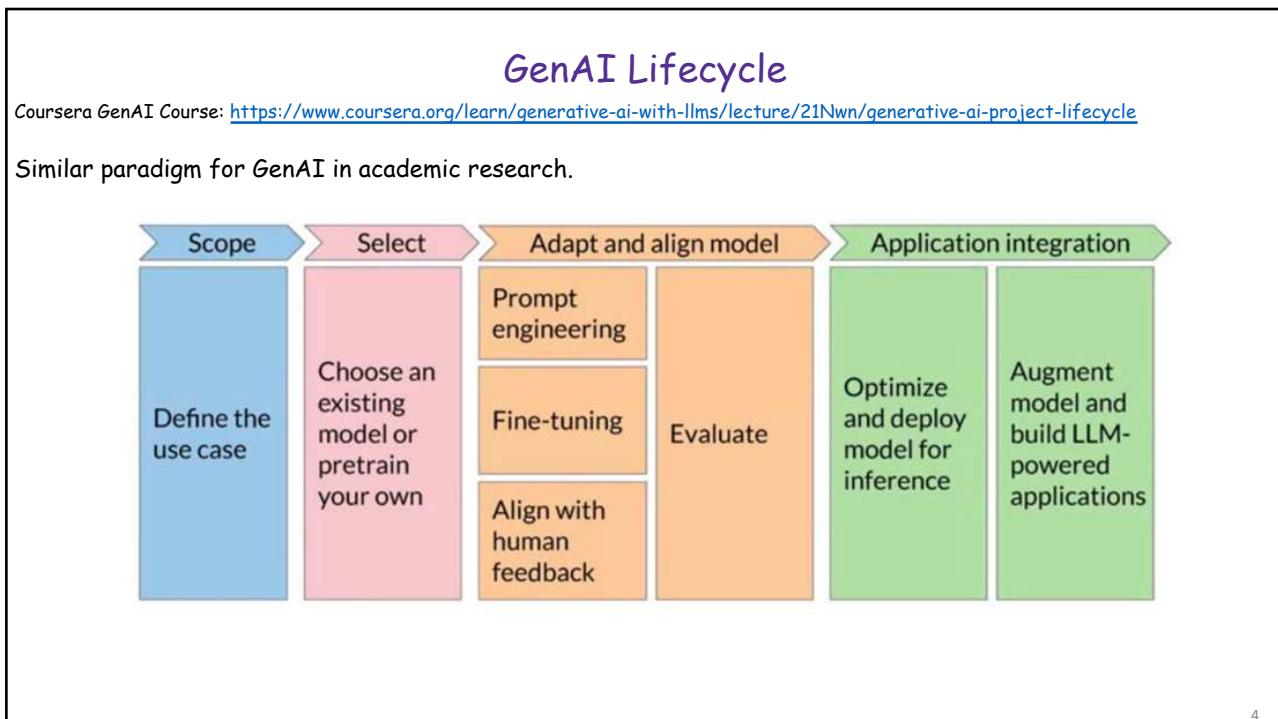
2

1



3

3



4

4

## Why Do We Need Posttraining?

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>; Hyperbolic Base Model Inference: <https://app.hyperbolic.xyz>  
 Talk by Barret Zoph and John Schulman: <https://docs.google.com/presentation/d/11KWCKUORnPpVMSY6vXqBeFSWo7fJcu6Q9yuR6vC1pzE>

Llama-3.1-405B-BASE

Demo API

Explain moon landing to a 6-year old

How would you explain the moon landing to a 6-year old? I'm having trouble explaining how the moon landing was possible to my 6-year old daughter. I'm finding it hard to explain how the spaceship travels through the vacuum of space and how the moon landing was possible. She's inquisitive and I'm not a scientist. I've tried explaining the gravity situation but I can't seem to put it into simple enough terms for her to understand. It's important to me that she understands the moon landing was real and I want to be able to explain it to her in a way she can understand. Any tips on how to do this? We talked about the moon landing and she asked how it was possible. I said it was because of gravity and she asked what gravity was. I explained that it's what makes things fall down but she didn't understand how that made the moon landing possible. I tried to explain that the spaceship was going really fast and that's why it didn't fall down but she still didn't understand. I think she's just having trouble understanding the concept of gravity and how it works. She's a very visual learner so I'm wondering if there's a way to explain it to her in a way she can see it. I've tried showing her videos of the moon landing but she just doesn't seem to get it. I think she needs a more hands-on explanation. I'm not sure how to explain it to her in a way she can understand. Any suggestions would be greatly appreciated. Thank you! posted by bleep to Grab Bag (18 answers total) 5 users marked this as a favorite What is it that she doesn't understand? Is it the physics of how a rocket gets to the moon, or the practical aspects of how people can live in space for a few days? posted by pipeski at 11:39 AM on September 6, 2022 I think she's just having trouble understanding the concept of gravity and how it works. I think you need to start there. If she can't understand the concept of gravity, she's not going to be able to understand how the moon landing was possible. posted by NotMyselfRightNow at 11:42 AM on September 6, 2022 [6 favorites] If you want to explain how a rocket works, I would use a balloon. Blow up a balloon and let go. Explain that a rocket is like a balloon that can keep blowing itself up. posted by pipeski at

Enter text here and AI will help complete the sentence

5

## Why Do We Need Posttraining?

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>; Hyperbolic Base Model Inference: <https://app.hyperbolic.xyz>  
 Talk by Barret Zoph and John Schulman: <https://docs.google.com/presentation/d/11KWCKUORnPpVMSY6vXqBeFSWo7fJcu6Q9yuR6vC1pzE>

Llama-3.1-405B

Demo API

Explain moon landing to a 6-year old

Oh boy, are you ready for an amazing story?

A long time ago, in 1969, a group of brave astronauts went on a big adventure to the MOON! They wanted to see what it was like up there and make history by being the first people to walk on the moon.

Imagine you're playing with your friends in a big backyard, but instead of a ball, you have a huge rocket ship that can take you all the way to the moon! That's what these astronauts did. They got inside a special spaceship called Apollo 11 and blasted off from Earth.

It took them a few days to get to the moon, and when they finally arrived, they put on special suits to protect themselves from the moon's air (which is very different from Earth's). Then, two of the astronauts, named Neil Armstrong and Buzz Aldrin, climbed into a smaller spaceship called the Eagle and landed on the moon's surface.

Can you imagine taking your first steps on the moon? It must have been so cool! Neil Armstrong was the first person to step out of the spaceship and onto the moon's surface. He said, "That's one small step for man, one giant leap for mankind!" which means it was a big deal for humans to achieve this amazing feat.

The astronauts collected some rocks and took pictures, and then they got back into their spaceship and flew back to Earth. They were heroes and everyone was so proud of them!

Isn't that an awesome story? Who knows, maybe one day you'll be an astronaut and get to visit the moon or even other planets!

[View API](#)

6

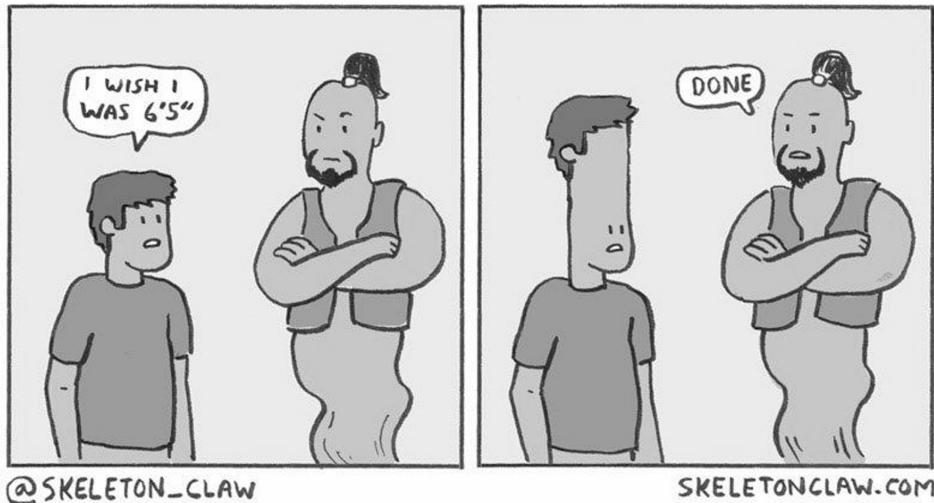
6

3

## AI Misalignment

- **Misalignment:** AI behaves in a way humans do not want.

GENIE 2

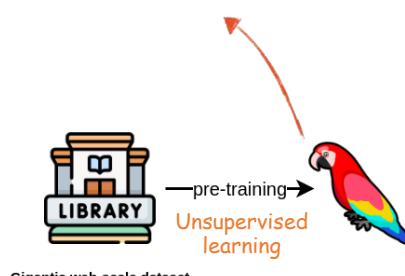


7

7

## Pretraining

Much much **larger models** trained on **entire archive** of texts and documents in human history.



## Posttraining



- **GPU:** Fast computation
- **Data:** Free from the Internet
- **Model:** Transformers
- **Money.....**

Training Costs	Pre-Training	Context Extension	Post-Training	Total
in H800 GPU Hours in USD	2664K \$5.328M	119K \$0.238M	5K \$0.01M	2788K \$5.576M

Table 1 | Training costs of DeepSeek-V3, assuming the rental price of H800 is \$2 per GPU hour.



Address the **alignment** and **safety** issues.

Slightly adjust the pre-trained model for **subsequent tasks**.

Training language models to follow instructions with human feedback  
L.Qiuyang, J.Wu, X.Jiang, D.Almeida ... - Advances in ..., 2022 - proceedings.neurips.cc  
Making language models bigger does not inherently make them better at following a user's intent. For example, large language models can generate outputs that are untruthful, toxic, or simply not helpful to the user. In other words, these models are not aligned with their users. In this paper, we show an avenue for aligning language models with user intent on a wide range of tasks by fine-tuning with human feedback. Starting with a set of labeler-written prompts and prompts submitted through a language model API, we collect a dataset of ...  
☆ Save 99 Cite Cited by 4049 Related articles All 10 versions ⚡

8

Andrej Karpathy's Deep Dive into LLM: [www.youtube.com/watch?v=7xTGNNLPyMI](https://www.youtube.com/watch?v=7xTGNNLPyMI)

8

## Posttraining vs. Pretraining

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>

Talk by Barret Zoph and John Schulman: <https://docs.google.com/presentation/d/11KWCKUORnPpVMSY6vXqBeFSWo7fJcu6Q9yuR6vC1pzE>

- Much **less compute** (so, much **cheaper** as well) than pretraining.
- Uses SFT and RLHF to **align** the models with **human preferences**.
- Teaches the model how to **use tools**.
  - Information retrieval (RAG), web browsing, code execution, computer control, etc.
- Crafts the model **personality**.
- Sets **refusal/safety** behavior.
  - "As an AI Language Model....."
- The effect of posttraining heavily relies on the **capability of the pretrained base model**.

9

## Mis-misalignment



Over-alignment also makes LLMs dumber.....

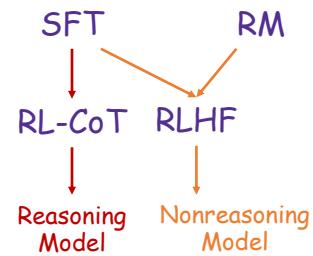
10

10

## Posttraining Components

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>  
 Talk by Barret Zoph and John Schulman: <https://docs.google.com/presentation/d/11KWCKUORnPpVMSY6vXqBeFSWo7fJcu6Q9yuR6vC1pzE>

- Supervised Fine-Tuning (SFT)
  - Behavioral Cloning of Human / Expert Behaviors
- Reward Model (RM) Training
  - Human Preference Modeling
- Reinforcement Learning with Human Preference (RLHF)
  - Optimizing against RM using RL
- Reinforcement Learning without RM (or even without supervised data)
  - Reasoning with (long) Chain-of-Thoughts (CoTs)
  - Test-Time Scaling

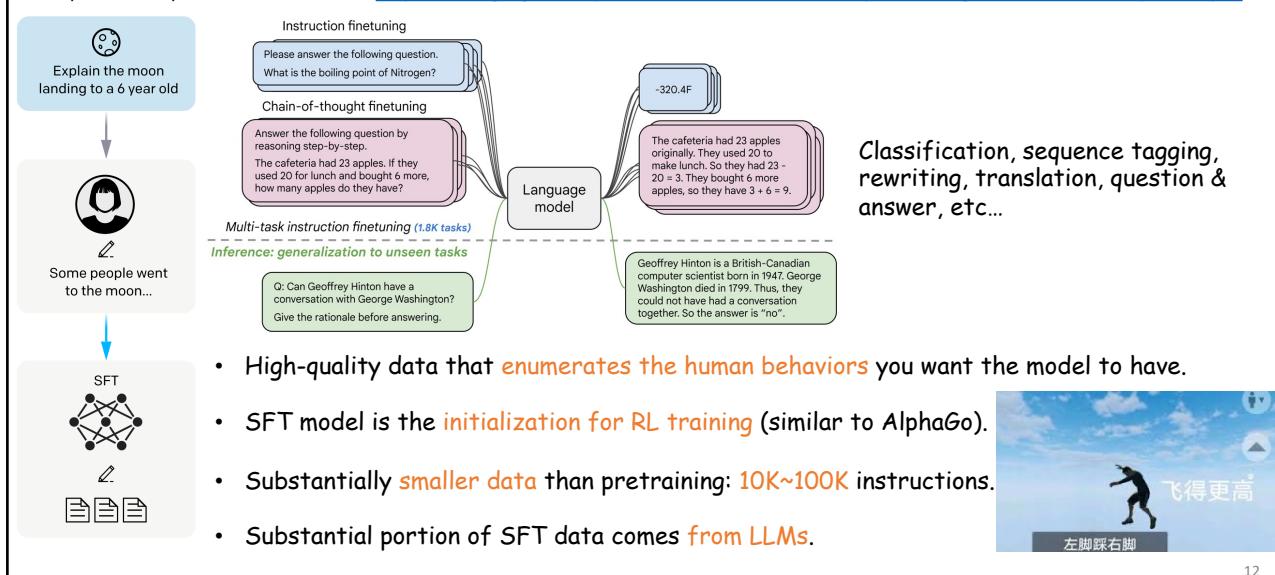


11

11

## Supervised Fine-Tuning

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>  
 Talk by Barret Zoph and John Schulman: <https://docs.google.com/presentation/d/11KWCKUORnPpVMSY6vXqBeFSWo7fJcu6Q9yuR6vC1pzE>



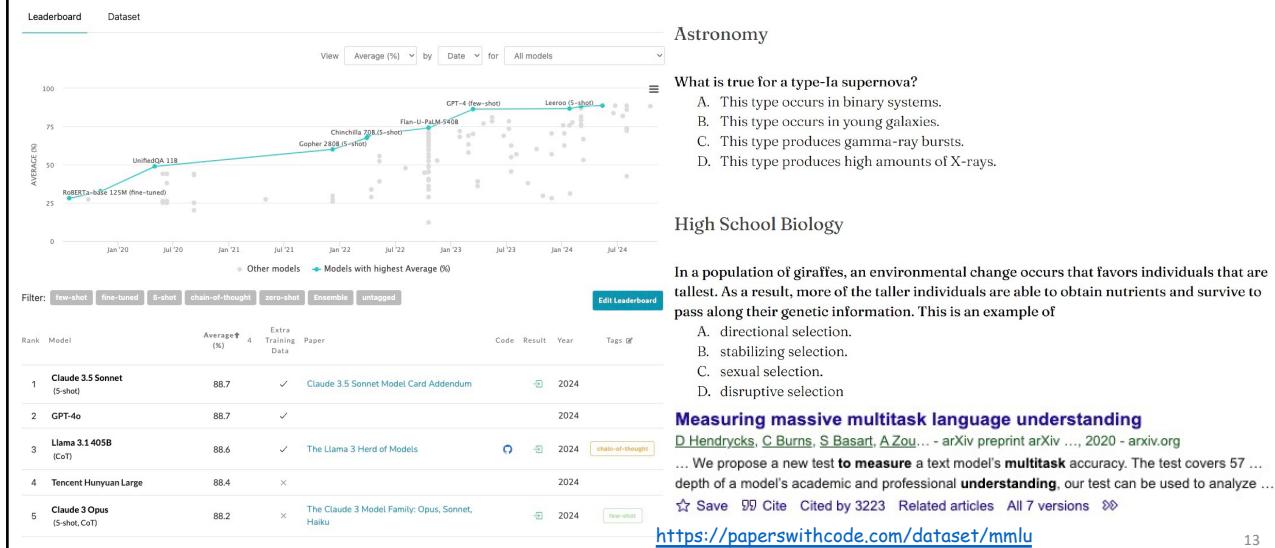
12

12

## Massive Multitask Language Understanding (MMLU)

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>; MMLU Paper: <https://arxiv.org/pdf/2009.03300>

- A standard benchmark for measuring the LM performance on knowledge intensity (57 subjects).



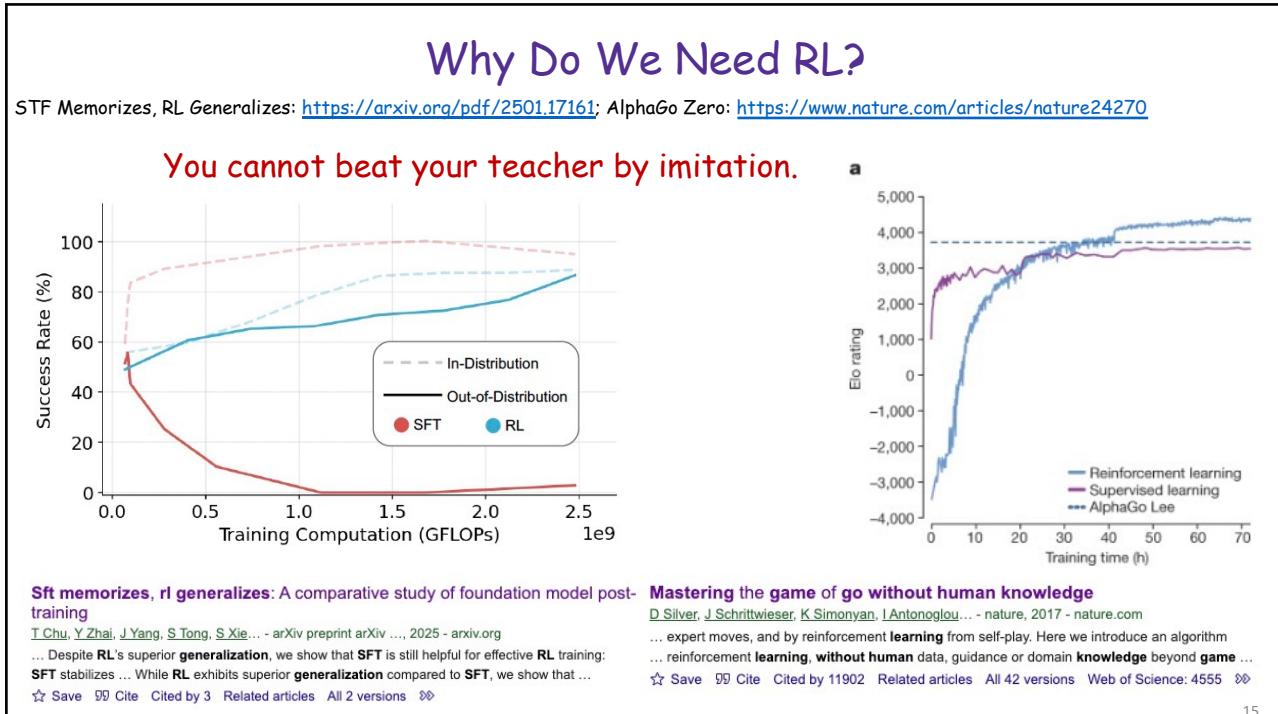
13

## Agenda

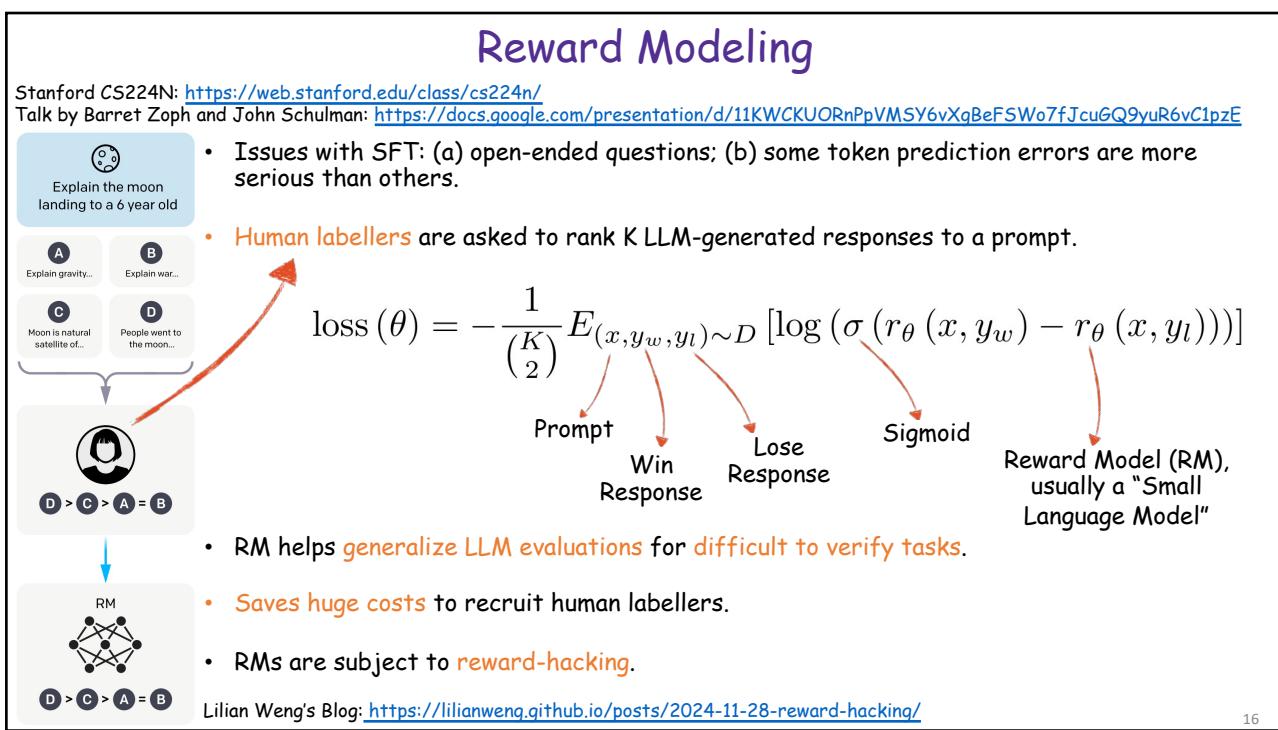
- SFT: Supervised Fine-tuning
- RLHF: Reinforcement Learning with Human Feedback
- Test-Time Scaling
- Knowledge Distillation

14

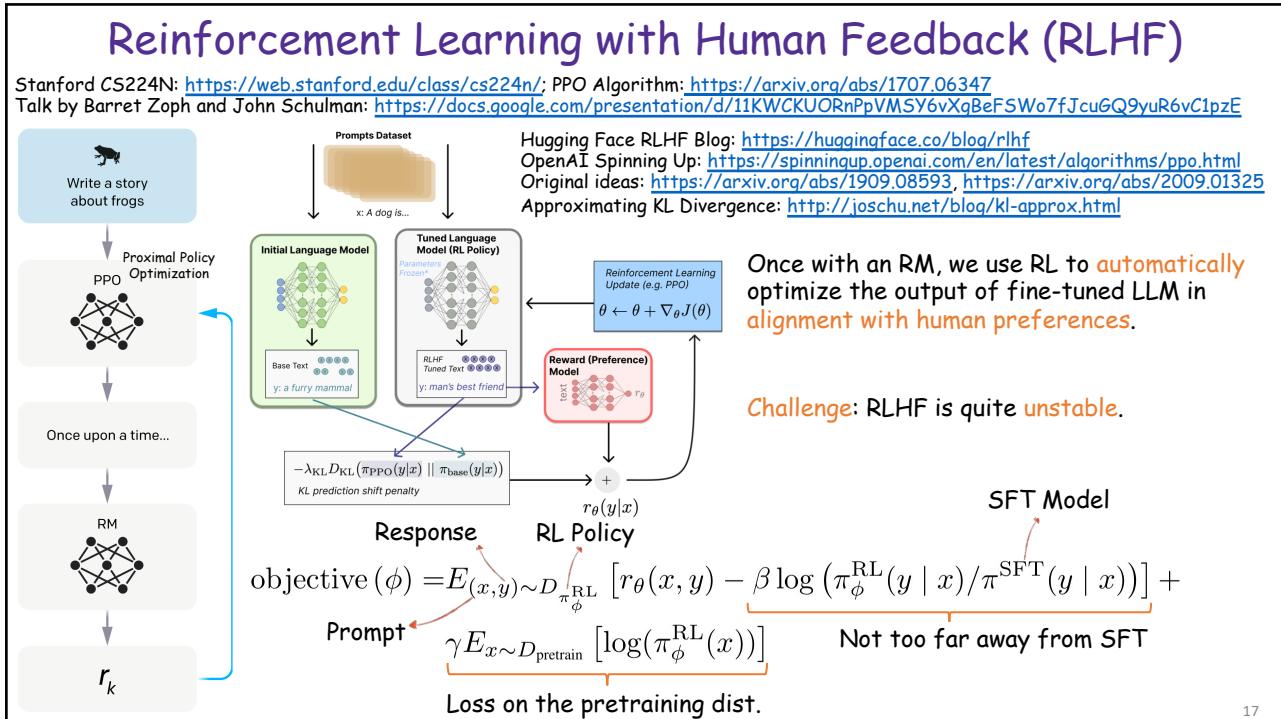
14



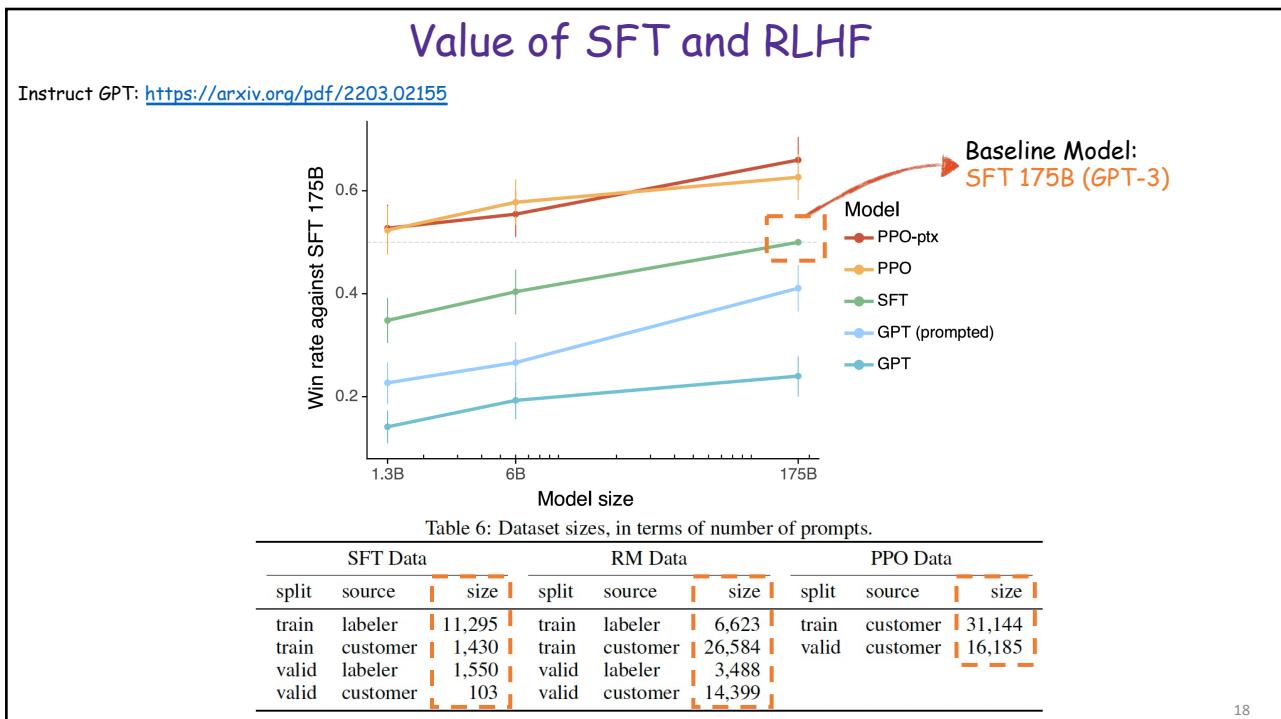
15



16



17



18

## Limitations with RL + RM

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>; OpenAI Reward Hacking: [openai.com/blog/faulty-reward-functions/](https://openai.com/blog/faulty-reward-functions/)  
Talk by Barret Zoph and John Schulman: <https://docs.google.com/presentation/d/11KWCKUORnPpVMSY6vXaBeFSWo7fJcu6Q9yuR6vC1pzE>

- Reward hacking is a common issue in RL.

TECHNOLOGY

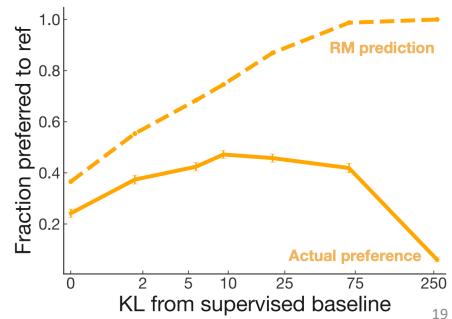
Google shares drop \$100 billion after its new AI chatbot makes a mistake

February 9, 2023 · 10:15 AM ET

- Human preferences are unreliable, so the LLM are rewarded to produce responses that seem authoritative and helpful, regardless of truth: Make-up facts and hallucinations.

- Models of human preferences are even more unreliable.

Learning to summarize from human feedback: <https://arxiv.org/pdf/2009.01325.pdf>

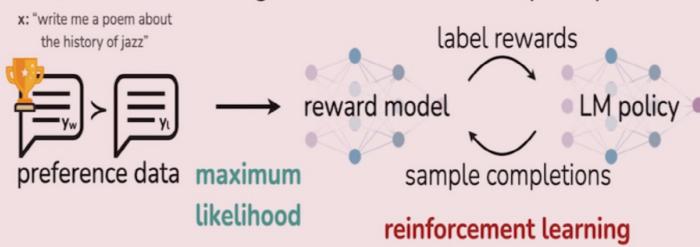


19

## Direct Preference Optimization (DPO)

Stanford CS224N: <https://web.stanford.edu/class/cs224n/>  
Talk by Barret Zoph and John Schulman: <https://docs.google.com/presentation/d/11KWCKUORnPpVMSY6vXaBeFSWo7fJcu6Q9yuR6vC1pzE>

### Reinforcement Learning from Human Feedback (RLHF)



### Direct Preference Optimization (DPO)



- RL is unstable and challenging to implement.
- Open-source (non-reasoning) LLMs ([https://huggingface.co/spaces/open-llm-leaderboard/open\\_llm\\_leaderboard](https://huggingface.co/spaces/open-llm-leaderboard/open_llm_leaderboard)) mostly use DPO.

$$\text{DPO-Loss} = -\mathbb{E}_{(x, y_w, y_l) \sim D} \left[ \log \sigma(\beta \log \frac{p_\theta^{RL}(y_w|x)}{p_{PT}^{RL}(y_w|x)} - \beta \log \frac{p_\theta^{RL}(y_l|x)}{p_{PT}^{RL}(y_l|x)}) \right]$$

From Human Rankings

Reward for winning sample

Reward for losing sample

DPO Paper: <https://arxiv.org/pdf/2305.18290.pdf>

20

20

## Hallucination

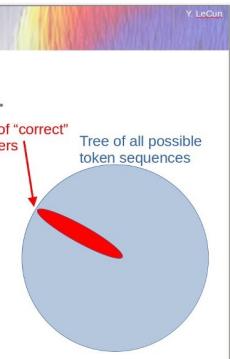
### The ChatGPT Lawyer Explains Himself

In a cringe-inducing court hearing, a lawyer who relied on A.I. to craft a motion full of made-up case law said he “did not comprehend” that the chat bot could lead him astray.



Steven A. Schwartz told a judge considering sanctions that the episode had been “deeply embarrassing.” Jefferson Siegel for The New York Times

**Unpopular Opinion about AR-LLMs**



- ▶ Auto-Regressive LLMs are **doomed**.
- ▶ They cannot be made factual, non-toxic, etc.
- ▶ They are not controllable
- ▶ Probability  $e$  that any produced token takes us outside of the set of correct answers
- ▶ Probability that answer of length  $n$  is correct:  

$$P(\text{correct}) = (1-e)^n$$
- ▶ **This diverges exponentially.**
- ▶ **It's not fixable (without a major redesign).**

- **Hallucination:** LLM learns the format but the content. Let alone the rationale and insights.

问一下知乎的医生们，以后的患者要是用DeepSeek问答的内容和你们对线要怎么办？

关注问题 写回答 邀请回答 好问题 2 1条评论 分享 ...

21

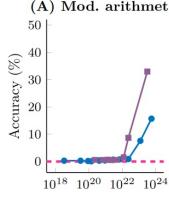
21

## Emergent Abilities

- **Emergent Abilities:** An ability not present in smaller models but present in larger models.
- **Phase-change in physics:** Quantitative changes in the system result in qualitative changes in behavior.

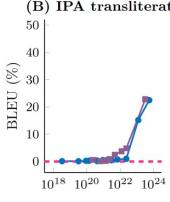


(A) Mod. arithmetic



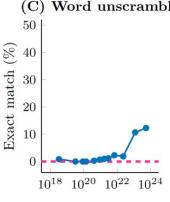
Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>18</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>20</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>22</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>24</sup>	~15%	~35%	~10%	~10%	~15%	~1%

(B) IPA transliterate



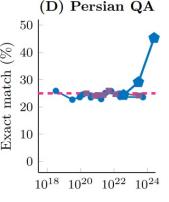
Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>18</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>20</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>22</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>24</sup>	~15%	~25%	~10%	~10%	~20%	~1%

(C) Word unscramble

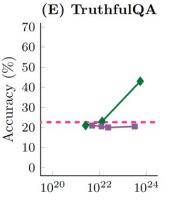


Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>18</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>20</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>22</sup>	~1%	~1%	~1%	~1%	~1%	~1%
10 <sup>24</sup>	~15%	~25%	~10%	~10%	~20%	~1%

(D) Persian QA

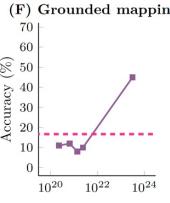


Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>18</sup>	~25%	~25%	~25%	~25%	~25%	~25%
10 <sup>20</sup>	~25%	~25%	~25%	~25%	~25%	~25%
10 <sup>22</sup>	~25%	~25%	~25%	~25%	~25%	~25%
10 <sup>24</sup>	~45%	~45%	~45%	~45%	~45%	~25%



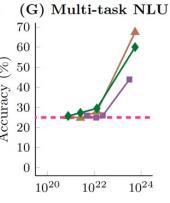
Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>20</sup>	~20%	~20%	~20%	~20%	~20%	~20%
10 <sup>22</sup>	~20%	~20%	~20%	~20%	~20%	~20%
10 <sup>24</sup>	~45%	~45%	~45%	~45%	~45%	~20%

(F) Grounded mappings



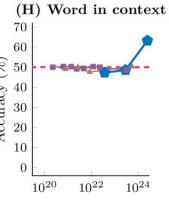
Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>20</sup>	~15%	~15%	~15%	~15%	~15%	~15%
10 <sup>22</sup>	~15%	~15%	~15%	~15%	~15%	~15%
10 <sup>24</sup>	~65%	~65%	~65%	~65%	~65%	~15%

(G) Multi-task NLU



Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>20</sup>	~25%	~25%	~25%	~25%	~25%	~25%
10 <sup>22</sup>	~25%	~25%	~25%	~25%	~25%	~25%
10 <sup>24</sup>	~65%	~65%	~65%	~65%	~65%	~25%

(H) Word in context



Model scale (FLOPs)	LaMDA	GPT-3	Gopher	Chinchilla	PaLM	Random
10 <sup>20</sup>	~50%	~50%	~50%	~50%	~50%	~50%
10 <sup>22</sup>	~50%	~50%	~50%	~50%	~50%	~50%
10 <sup>24</sup>	~65%	~65%	~65%	~65%	~65%	~50%

**Emergent abilities of large language models**  
J Wei, Y Tay, R Bommasani, C Raffel, B Zoph... - arXiv preprint arXiv ..., 2022 - arxiv.org  
... an ability to be emergent if it is not present in smaller models but is present in larger models. ... We have discussed emergent abilities of language models, for which meaningful ...  
☆ Save 99 Cite Cited by 1256 Related articles All 11 versions

22

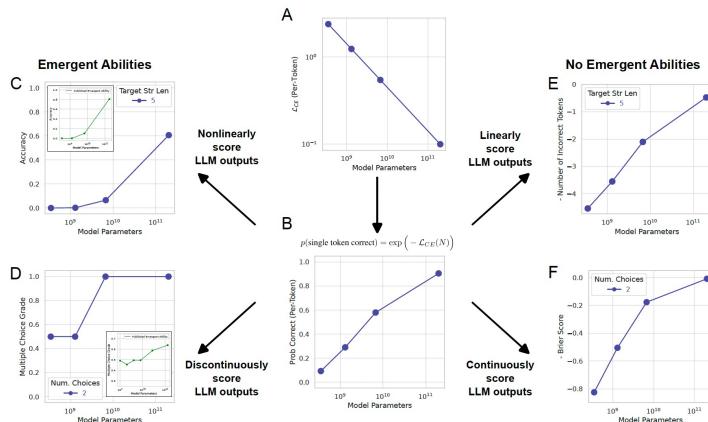
22

## Are Emergent Abilities of LLMs a Mirage?

$$\text{Multiple Choice Grade} \stackrel{\text{def}}{=} \begin{cases} 1 & \text{if highest probability mass on correct option} \\ 0 & \text{otherwise} \end{cases}$$

$$\text{Exact String Match} \stackrel{\text{def}}{=} \begin{cases} 1 & \text{if output string exactly matches target string} \\ 0 & \text{otherwise} \end{cases}$$

- **Emergent Abilities** may be attributed to the choice of nonlinear or discontinuous metrics, whereas linear or continuous metrics produce smooth performance changes.



NeurIPS 2023 Outstanding Main Track Paper

### Power-law of Scaling

$$\mathcal{L}_{CE}(N) = \left(\frac{N}{c}\right)^\alpha$$

$$p(\text{single token correct}) = \exp(-\mathcal{L}_{CE}(N)) = \exp(-(N/c)^\alpha)$$

$$\text{Accuracy}(N) \approx p_N(\text{single token correct})^{\text{num. of tokens}} = \exp(-(N/c)^\alpha)^L$$

Figure C

Figure E

$$\text{Token Edit Distance}(N) \approx L(1 - p_N(\text{single token correct})) = L(1 - \exp(-(N/c)^\alpha))$$

Are emergent abilities of large language models a mirage?

R Schaeffer, B Miranda, ... Advances in Neural ..., 2024 - proceedings.neurips.cc

... be interpreted as claiming that **large language models** cannot display **emergent abilities**:

rather, our message is that some previously claimed **emergent abilities** appear to be mirages ...

☆ Save 99 Cite Cited by 129 Related articles All 9 versions

23

23

## Agenda

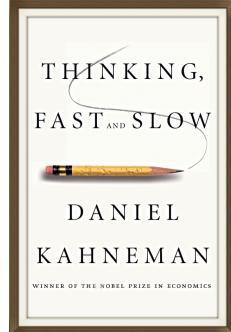
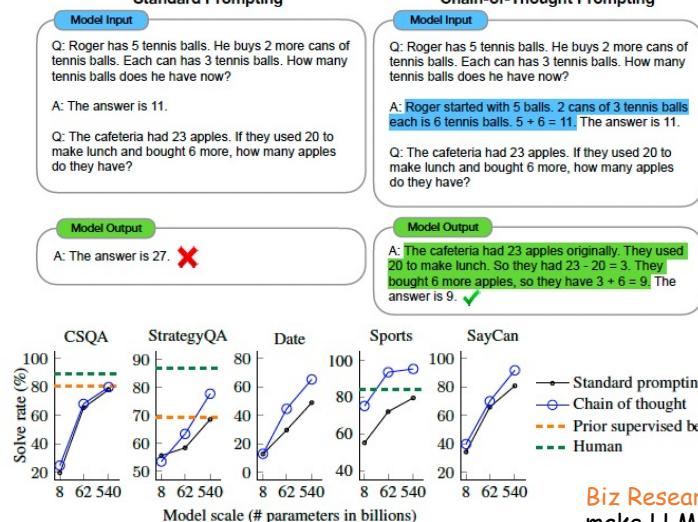
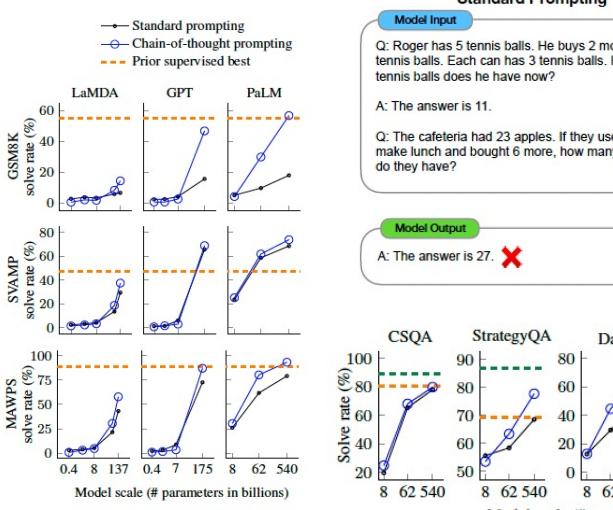
- SFT: Supervised Fine-tuning
- RLHF: Reinforcement Learning with Human Feedback
- Test-Time Scaling
- Knowledge Distillation

24

24

## Chain-of-Thought (CoT)

- A series of **intermediate reasoning steps** significantly improves the ability of large language models for complex reasoning.



**Pros:** Reasoning,  
Debug

**Cons:** Cost, Speed

**Biz Research:** CoT may make LLMs **overly rational**.

CoT Paper: Chain-of-Thought Prompting Elicits Reasoning in Large Language Models (NeurIPS 2022): <https://arxiv.org/pdf/2201.11903.pdf>

25

# Tree-of-Thought (ToT)

- Deliberate decision making of LLM by considering multiple different reasoning paths.

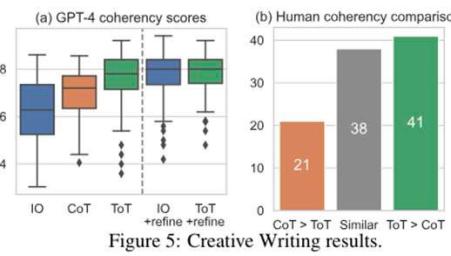
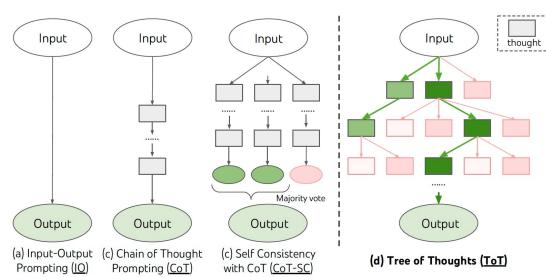


Figure 5: Creative Writing results.

Tree of thoughts: Deliberate problem solving with large language models

[S Yao, D Yu, J Zhao, I Shafran...](#) - Advances in neural ... , 2023 - proceedings.neurips.cc

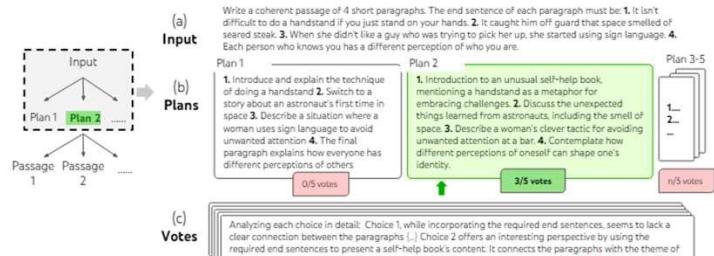
Abstract Language models are increasingly being deployed for general problem solving

across a wide range of tasks, but are still confined to token-level, left-to-right decision-

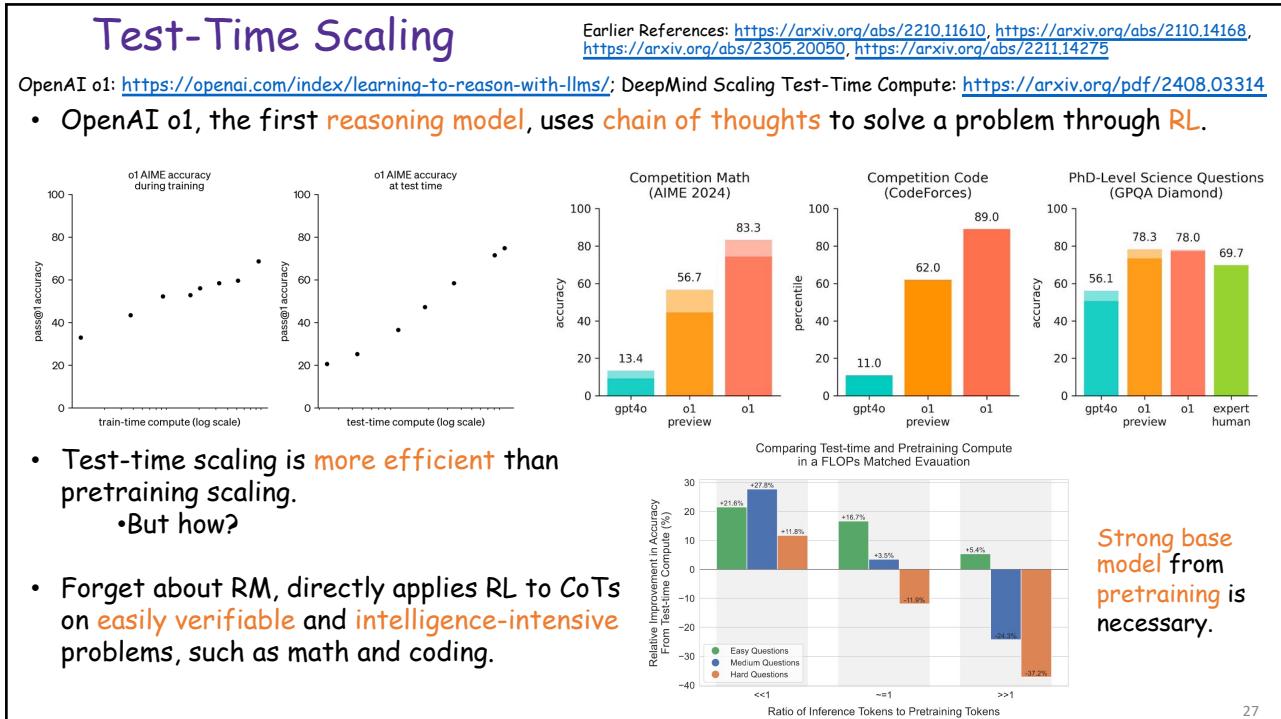
making processes during inference. This means they can fall short in tasks that require them to make inferences based on incomplete or ambiguous information.

strategic lookahead, or where initial decisions play a pivotal role. In this paper, we introduce a new framework for language model inference that addresses these challenges.

- Thought decomposition
  - Thought generator
  - State evaluator
  - Search algorithm

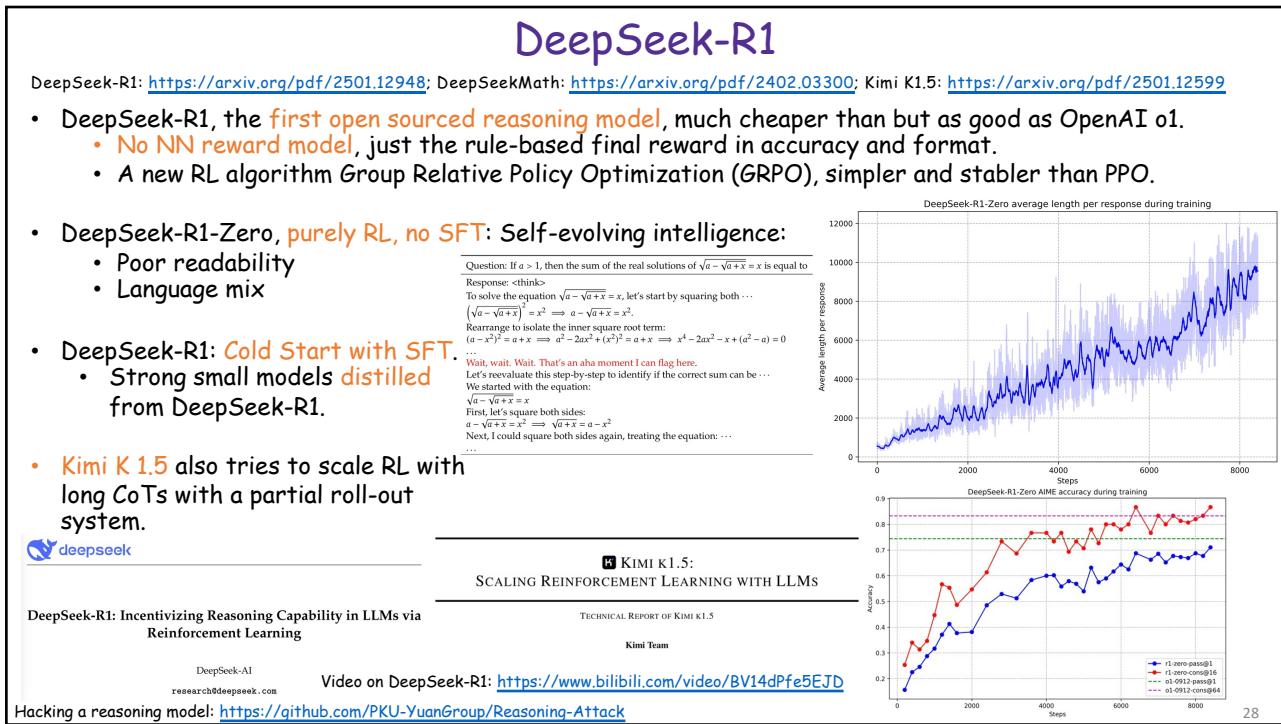


26



27

27



28

28

JOURNAL ARTICLE CORRECTED PROOF  
**Generative AI at Work** 

Erik Brynjolfsson, Danielle Li, Lindsey Raymond  
*The Quarterly Journal of Economics*, qjae044, <https://doi.org/10.1093/qje/qjae044>  
 Published: 04 February 2025 Article history ▾

PDF Split View Cite Permissions Share ▾

**Abstract**

We study the staggered introduction of a generative AI-based conversational assistant using data from 5,172 customer-support agents. Access to AI assistance increases worker productivity, as measured by issues resolved per hour, by 15% on average, with substantial heterogeneity across workers. The effects vary significantly across different agents. Less experienced and lower-skilled workers improve both the speed and quality of their output, while the most experienced and highest-skilled workers see small gains in speed and small declines in quality. We also find evidence that AI assistance facilitates worker learning and improves English fluency, particularly among international agents. While AI systems improve with more training data, we find that the gains from AI adoption are largest for moderately rare problems, where human agents have less baseline experience but the system still has adequate training data. Finally, we provide evidence that AI assistance improves the experience of work along several dimensions: customers are more polite and less likely to ask to speak to a manager.

**Generative AI at work**  
 E Brynjolfsson, D Li, L Raymond - *The Quarterly Journal of ...*, 2025 - academic.oup.com  
 We study the staggered introduction of a generative AI-based conversational assistant using data from 5,172 customer-support agents. Access to AI assistance increases worker ...  
 Save  Cite Cited by 1002 Related articles All 18 versions 

## Reasoning Models in Biz Research?

- Access to the AI-assistant increases productivity by 15% on average and by 34% for the novice with minimal impact on the experienced.
  - Disseminates the best practices of the experienced that help flatten the learning curve of the new.
  - AI reduces the marginal cost of distributing intelligence.
- 如何看待镇江部署DeepSeek，称「建成600台算力服务器集群，单日处理量为全市公务员十年工作量」？  
 2月19日，镇江举行新闻发布会，DeepSeek正式登陆镇江，完成本地化部署上线，这是镇江市推进数字经济高质量发展的关键举措。
- Economic impact of reasoning model deployment at large?
  - Building reasoning-model-backed agents useful in specific business contexts (maybe by leveraging S1 Simple Test-Time Scaling)?
  - Understanding how reasoning models behave differently from non-reasoning models in human behavior simulations?

<https://www.zhihu.com/question/12918439244>  
<https://arxiv.org/pdf/2501.19393.pdf>

29

## Agenda

- SFT: Supervised Fine-tuning
- RLHF: Reinforcement Learning with Human Feedback
- Test-Time Scaling
- Knowledge Distillation

30

30

## Knowledge Distillation (KD)

MIT 6.5940 Efficient DL Computing: [efficientml.ai](http://efficientml.ai)

- Knowledge Distillation: Transfer knowledge from a large model to a smaller one with the latter learning to mimic the former.

The diagram illustrates the process of Knowledge Distillation. An input  $x$  is processed by both a Teacher model (with layers 1, 2, ..., m) and a Student (distilled) model (with layers 1, 2, ..., n). The Teacher model's output is passed through a Softmax function at temperature  $T = t$  to produce soft labels. The Student model's output is also passed through a Softmax function at the same temperature  $T = t$  to produce soft predictions. Both sets of predictions are used to calculate the KL Divergence loss. Additionally, the Student model's output is passed through a Softmax function at  $T = 1$  to produce hard predictions, which are compared against a hard label  $y$  (ground truth) to calculate the student loss. The total loss is the weighted sum of the KL Divergence loss and the student loss.

Model	AIME 2024		MATH-500		GPQA		LiveCode Bench		CodeForces	
	pass@1	cons@64	pass@1	pass@1	Diamond	Bench	pass@1	rating	pass@1	1633
GPT-4o-0513	9.3	13.4	74.6	49.9	32.9	16.9	954			
Claude-3.5-Sonnet-1022	16.0	26.7	78.3	65.0	38.9	717				
OpenAI-01-mini	63.6	80.0	90.0	60.0	53.8	1820				
QwQ-32B-Preview	50.0	60.0	90.6	54.5	41.9	1316				
DeepSeek-R1-Distill-Qwen-1.5B	28.9	52.7	83.9	33.8	16.9	954				
DeepSeek-R1-Distill-Qwen-32B	55.5	83.3	92.8	49.1	37.6	1189				
DeepSeek-R1-Distill-Qwen-14B	69.7	80.0	93.9	50.1	53.1	1481				
DeepSeek-R1-Distill-Qwen-32B	72.6	83.3	94.3	62.1	57.2	1691				
DeepSeek-R1-Distill-Llama-8B	50.4	80.0	89.1	49.0	39.6	1205				
DeepSeek-R1-Distill-Llama-70B	70.0	86.7	94.5	65.2	57.5	1633				

Table 5 | Comparison of DeepSeek-R1 distilled models and other comparable models on reasoning-related benchmarks.

31

31

## KD for Market Research

- We have limited human data and a lot of LLM-generated data. How to correctly identify human preferences?
- Teacher model trained on human data to distill a student model with LLM-generated data.

**ESTIMATION WITH AI-AUGMENTED DATA**

Step 1. Obtain an estimator  $\hat{\theta}$  to  $\theta^*$ , where  $P(y=j|x, z) = g_j(x, z; \theta^*)$ , using the primary data.

Step 2. With the auxiliary data, we construct the estimator  $\hat{\theta}^{AAE}$  as **AI-Augmented Estimator**

$$\text{Distillation: } \hat{\theta}^{AAE} \in \arg \max_{\beta \in \mathbb{R}^d} \left\{ \hat{Q}(\hat{\theta}; \beta) = \frac{1}{n} \sum_{i=1}^n \sum_{j \in \mathcal{K}^+} g_j(x_i, z_i; \hat{\theta}) \log \sigma_j(x; \beta) \right\}.$$

The graph plots Accuracy (%) on the y-axis (ranging from 50 to 80) against Cost (\$1K) on the x-axis (ranging from 0.5 to 3.5). Three data series are shown: Traditional (orange diamonds), AAE (blue diamonds), and Naive (green diamonds). The AAE method consistently achieves the highest accuracy across all costs, while the Traditional method has the highest variance and the Naive method has the lowest accuracy.

**Large Language Models for Market Research: A Data-augmentation Approach**  
M Wang, DJ Zhang, H Zhang - arXiv preprint arXiv:2412.19363, 2024 - arxiv.org  
... our context, we present a **data-augmentation** statistical approach for extracting value from LLM-... **data augmentation approach** that allows us to use the AI-generated data to fit the **model**. ...  
☆ 保存 ⚡ 引用 被引用次数: 1 相关文章 所有 3 个版本 ⚡  
<https://arxiv.org/pdf/2412.19363.pdf>

**Large Language Models for Market Research: A Data-augmentation Approach**  
Mengxin Wang (Naveen Jindal School of Management, The University of Texas at Dallas), Dennis J. Zhang (Olin School of Business, Washington University in St. Louis), Heng Zhang (W. P. Carey School of Business, Arizona State University)

Large Language Models (LLMs) have transformed artificial intelligence by excelling in complex natural language processing tasks. Their ability to generate human-like text has opened new possibilities for market research, particularly in conjoint analysis, where understanding consumer preferences is essential but often resource-intensive. Traditional survey-based methods face limitations in scalability and cost, making LLM-generated data a promising alternative. However, while LLMs have the potential to simulate real consumer behavior, recent studies highlight a significant gap between LLM-generated and human data, with biases introduced when substituting between the two. In this paper, we address this gap by proposing a novel statistical data augmentation approach that efficiently integrates LLM-generated data with real data in conjoint analysis. Our method leverages transfer learning principles to debias the LLM-generated data using a small amount of human data. This results in statistically robust estimators with consistent and asymptotically normal properties, in contrast to naive approaches that simply substitute human data with LLM-generated data, which can exacerbate bias. We validate our framework through an empirical study on COVID-19 vaccine preferences, demonstrating its superior ability to reduce estimation error and save data and costs by 24.9% to 79.8%. In contrast, naive approaches fail to save data due to the inherent biases in LLM-generated data compared to human data. Another empirical study on sports car choices validates the robustness of our results. Our findings suggest that while LLM-generated data is not a direct substitute for human responses, it can serve as a valuable complement when used within a robust statistical framework.

Subjects: Artificial Intelligence (cs.AI); Machine Learning (cs.LG); Methodology (stat.ME); Machine Learning (stat.ML)

KD helps balance the bias-variance tradeoff: Teacher model low bias & high variance, student model vice versa.

32

32